

A GoLite to JavaScript compiler

Guillaume Labranche

William Bain

Si Mei Zhang

April 14, 2015

Contents

1	Introduction	1
2	Lexing	1
2.1	Tokens	1
2.2	Visualizing tokenization	2
2.3	Semicolon insertion	2
3	Parsing	2
4	Weeding	2
5	Type checking	2
5.1	Typechecker architecture	2
5.2	Symbol table	3
5.3	Type and symbol class hierarchy	3
6	Pretty printers	4
7	Code generation	4
7.1	Output considerations: JavaScript	4
7.2	Generation techniques	4

1 Introduction

We initially considered C, Java, and Python as potential implementation languages, based on our shared familiarity with them. We decided not to use C because of the extra development costs incurred by using a language without features like memory safety. We opted for Java over Python because we judged there to be more resources for compiler development available.

We decided to use SableCC 3 because it made it easy to implement a full scanner/parser toolchain with utilities for AST traversal, and because there was in-class support for it. This was very useful, for instance, in the case of semicolon insertion.

TODO target language choice

2 Lexing

2.1 Tokens

We had 6 categories of tokens.

1. Text literals: runes, raw strings, interpreted strings.

2. Numeric literals: integers of different forms and floats.
3. Keywords: simply all the words that cannot be identifiers.
4. Operators and delimiters: simply a list of all of them.
5. Identifiers: letter, underscore and digits, with the latter not allowed as first character.
6. Comments: single line and multi-line.

Initially we scanned text literals entirely in the lexer, creating tokens which obscured the internal structure of the literals—in particular, the presence of escape characters. This was fine while we were doing the parsing and type checking, but we resolved to improve on it before turning to code generation. We opted to introduce new scanning modes which preserved whitespace and read escape sequences as separate tokens, and then implemented the literals as a whole in the parser, so that their structures were preserved in the AST.

2.2 Visualizing tokenization

In order to test the correctness of our scanning, we needed a way to display the tokenization on top of the source without altering its layout. A simple solution was to generate HTML from the stream of tokens, adding the token's name as a `title` attribute, and alternating colors to clearly see the break up of characters. This test is available with the `-dumptoks` flag.

TODO: insert graphic of highlighted source

2.3 Semicolon insertion

The fact that semicolons are optional makes it harder to construct our CST (TODO: Will, expand on that?). Therefore we must add semicolons to the tokenization following simple rules.¹ To do this with SableCC, we subclassed `Lexer` as `GoLexer` and implemented the `filter` method.

3 Parsing

TODO

4 Weeding

TODO

5 Type checking

5.1 Typechecker architecture

In accordance with the instructions our type checking is executed in a single pass without forward declaration.² We are working in Java and we implemented the typechecker as a subclass of the `DepthFirstAdapter` class provided by SableCC. This made it easy to traverse the AST without extra boilerplate. Since we did not have a reliable way of extending the AST classes generated by SableCC, we stored type information for the AST in a hash table mapping from nodes to GoLite types.³

In most cases it was sufficient to apply typechecks after the typechecker had recursed over the child nodes and typechecked them. There were two cases where more fine-grained control was necessary:

¹See https://golang.org/doc/effective_go.html#semicolons.

²See <https://mycourses2.mcgill.ca/d2l/1e/161312/discussions/threads/258597/View>.

³By extending I mean adding methods and properties to a class used by the parser (here probably `PExp`, the abstract expression production class), not just creating a subclass.

- In short assignment statements, the list of variables is implemented as a list of expressions for reasons having to do with the parser implementation. It was therefore necessary to stop the typechecker from typechecking the variables before they had been added to the symbol table.
- To typecheck function declarations and if statements, it was necessary to open a scope after having typechecked some but not all of the child nodes.

In each of these special cases we overrode the case method which controls the recursion over the child nodes.

A further implementation detail that is worth noting is our implementation of struct type declarations. To prevent code duplication, we treat struct fields as if they were variable declarations: we open a scope in the symbol table upon entering the struct and enter each field as if it were a variable. Then we pop the scope and use it to build the struct class.

5.2 Symbol table

Initially the symbol table was simply a scope with a reference to the parent scope. Traversing up and down the scopes meant overwriting the variable used to reference the scope, so we decided to make `SymbolTable` be what its name suggests, and take care of scoping up and down. It holds a double-ended queue (Java's `ArrayDeque`) to store the scopes. We have methods for different use cases: searching for an identifier through all scope levels, and searching only in the current scope.

We have also added loggers to enable the `-dumpsymtab` and `-dumpsymtaball` CLI options.

5.3 Type and symbol class hierarchy

We put considerable effort into developing and revising the hierarchy representing GoLite types and symbols (figure 1). We adopted the following goals, listed here roughly in descending order of precedence:

1. Type safety: as much as possible, it should not be possible to use an object in a place where it is not allowed. This should be detected at compile time.
2. Simplicity: there should not be more classes or objects than are necessary. For instance, we did not want to use a class `AliasTypeSymbol (Symbol s)` to wrap every type alias entered in the symbol table.
3. DRY: it should not be necessary to implement functionality in multiple places.

These considerations led us to a number of particular decisions. First, we place functions outside the type hierarchy. While in Go functions are first class citizens, with their type given by the function signature, in GoLite they are not. Therefore, instead of implementing functions as instances of `Variable`, which would potentially allow them to be used where another objects of another type are required, we implemented a `Function` class which represents a particular function and stores both the function's identifier and its associated type information directly.

Second, we made `Symbol` an interface which is implemented by the `Variable`, `Function`, and `AliasType` classes. The difficulty here was that since `AliasType` is a subclass of `Type`, it could not also inherit from an abstract `Symbol` class. By making `Symbol` an interface, we allow `AliasType` instances into the symbol table directly without thereby making it possible to enter other types—for instance the `BuiltInType` which we use for `int`, `bool`, `alias`, and `rune`—illicitly. The downside of this design is that the `getId` and `setId` methods required by `Symbol` had to be implemented at least twice. However, this was sufficiently trivial that it did not change our decision. In fact, the duplication is so trivial that we did not even create a shared superclass for `Function` and `Variable` to reduce the duplication there.⁴

A final decision worth noting is our use of the `VoidType` class. Since unfortunately Java does not make it possible to statically forbid null pointers, it is easy to introduce them accidentally. To help catch such cases, we use `VoidType` instead of `null` to represent the return type of a function that does not return a value.

⁴We *did*, however, create a `NamedType` class which is extended by both `BuiltInType` and `AliasType`. This class also implements the `isIdentical` and `getRepresentation` methods.

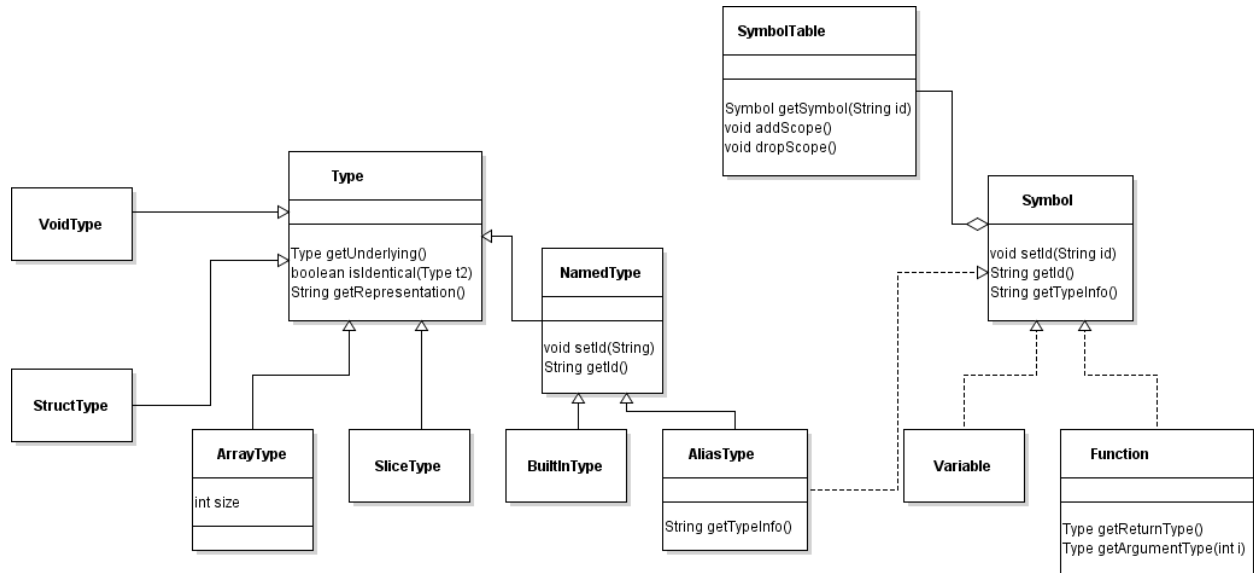


Figure 1: Class hierarchy for GoLite types and symbols (property and method listings are non-exhaustive)

6 Pretty printers

We extended the `PrettyPrinter` class we implemented in Milestone 1 to have it print extra type information. By overriding the `defaultOut` method, it was simple to have it look up and output the type of any node which is an instance of the `PExp` (expression production) class as an inline comment.

7 Code generation

7.1 Output considerations: JavaScript

TODO

7.2 Generation techniques

TODO